


Last week

- New constraint method where bias and threshold are learned through gradient descent
- Doesn't work




This week

- 1- Increase Firing rate constraint learning rate : 1x, 0.1x, 0.01x
- 2- New metrics for Lagrange and Two_losses
 - a) Gradient
 - b) Covariance matrices
- 3- Recording loss before and after adjusting FR
- 4- Toy model
- 5- Next steps

Increasing the LR of the FR constraint


KL

tag: terms/KL



MI

tag: terms/MI



What is the difference between KL and MI?

(Question from last week)

KL:

```
calculate_metrics(self, i, fitting_restriction) ...  
KL = self.logdet_numerator - self.logdet_denominator
```

MI:

```
writer.add_scalar(f"terms/MI", H_X - metrics.loss.mean().item(), iteration)  
loss=self.model.beta * KL,
```



Increasing LR for FR doesn't narrow the distribution of firing rates

histogram/r

1


histogram/r

10



histogram/r


100



Lagrange FR histogram
for comparison:

histogram/r

Lagrange




New Gradient metric: Lagrange vs Gamma

W_grad|
step 225 000

Lagrange


Mon Jan 22 2024 13:20:18 heure normale de l'Est nord-américain



W_grad
step 225 000

Two_losses

Mon Jan 22 2024 17:48:27 heure normale de l'Est nord-américain




Previous (Lagrange) covariance matrix: Doesn't look like it should

Reasons why:

John: Should be diagonal if we don't
multiply by G matrix

David: This is taken from 1 batch of 128

Maybe this is a noisy estimate




Both hypotheses seem to be wrong

Sampled from 1000 batches instead of just 1:

MI_numerator
step 1 000 000

Mon Jan 22 2024 16:06:12 heure normale de l'Est nord-américain


Lagrange



MI_denominator
step 1 000 000

Mon Jan 22 2024 16:06:12 heure normale de l'Est nord-américain


Lagrange



WCxW
step 1 000 000

Mon Jan 22 2024 16:06:12 heure normale de l'Est nord-américain

Lagrange




Ironically, these matrices look better with Gamma method

MI_numerator
step 1 000 000

Two_losses


Mon Jan 22 2024 20:38:27 heure normale de l'Est nord-américain



MI_denominator
step 1 000 000

Two_losses


Mon Jan 22 2024 20:38:27 heure normale de l'Est nord-américain



WCxW
step 1 000 000


Two_losses

Mon Jan 22 2024 20:38:27 heure normale de l'Est nord-américain



Why doesn't this work?

Toy problem with a 2D multivariate oriented gaussian



```

x1 = torch.tensor(2.0, requires_grad = True)
x2 = torch.tensor(1.0, requires_grad = True)

#y = torch.tensor(1.0, requires_grad = True)
cov = nn.Parameter(torch.tensor([[1,0.7], [0.7,1]]), requires_grad = False)
#xy = torch.tensor([x,y], requires_grad = True)
u = torch.tensor([1.0,1.5], requires_grad = False)

def forward(x):
    return torch.exp(-torch.matmul(torch.matmul(x-u,cov),x-u))

#z.requires_grad = True
optimizer = torch.optim.SGD([x1,x2], lr = 0.01)
optimizer2 = torch.optim.SGD([x2], lr = 0.2)
for i in range(1000):
    x = torch.stack((x1,x2))
    optimizer.zero_grad()
    optimizer2.zero_grad()
    #Let's start first optimization
    z = -forward(x)
    z.backward(retain_graph = True)
    optimizer.step()
    optimizer.zero_grad()

    #Lets start second optimization
    loss_2 = (x2 - torch.tensor(3))**2
    loss_2.backward()
    optimizer2.step()

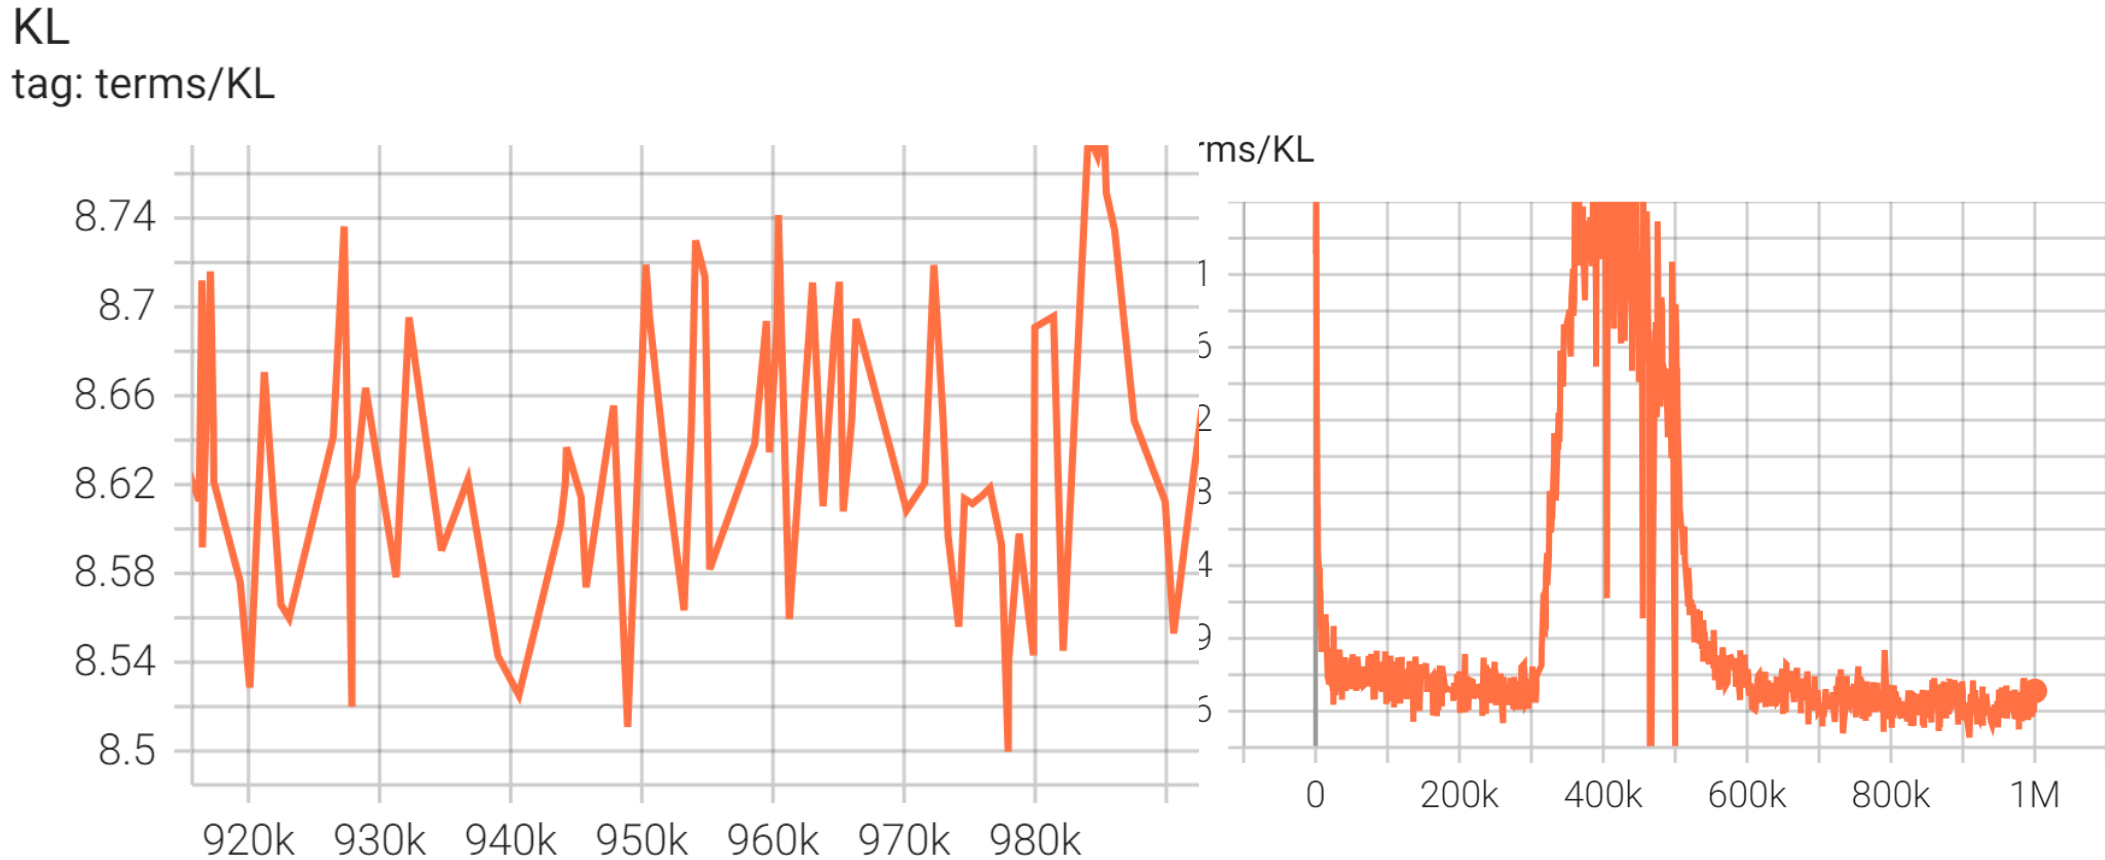
    optimizer2.zero_grad()
    print(x[0].item(), x[1].item(), z)

```

Results in a nutshell

- 1- I initially thought I could replicate the failure mode but turns out that was because I didn't let the algorithm run for long enough
- 2- This toy model works (As long as initial conditions aren't too far off from optimum)
- 3- I still don't know why the Gamma method doesn't, but that's okay

Recording KL before and after adjusting the firing rate



Ideas

- Back to 300x300 color model but with a learning rate that reduces over time (currently running)
- Figure out where algorithm takes most time to run. LU decomposition to optimize learning?
- In general, figure out how to make 300x300 model converge better